

Spark建置與開發實務

軟體發展組

莊家雋

Outline



- Part I : Quick tour of Spark
 - Why Spark
 - Prepare Lab VMs
 - Run first Spark example
 - {Java / Python / Scala } code via {Shell / Standalone application }
 - Introduction to FP, Java 8, and Scala
- Part II : Programming with RDD
 - RDD Concepts
 - RDD, RDD API, convert RDD
 - Load and save data
 - Spark application

Outline

- Part III: Run on Cluster
 - Cluster Manager
 - Standalone, YARN, Mesos
 - Components of Execution
 - Application, Job, stage, task, RDD, partition, DAG

Why Spark

- Compare with Hadoop ecosystem
 - More efficient execution
 - More unified program abstraction
 - More flexible program operation



transformation
<code>map(func)</code>
<code>filter(func)</code>
<code>flatMap(func)</code>
<code>sample(withReplacement, fraction, seed)</code>
<code>union(otherDataset)</code>
<code>distinct([numTasks])</code>
<code>groupByKey([numTasks])</code>
<code>reduceByKey(func, [numTasks])</code>
<code>sortByKey([ascending], [numTasks])</code>
<code>join(otherDataset, [numTasks])</code>
<code>cogroup(otherDataset, [numTasks])</code>
<code>cartesian(otherDataset)</code>

Exercise i-1:

Bring up Lab VMs

- Lab VM has
 - Ubuntu OS / Java8 / maven3 / spark 1.5 tarball / R
- Download and Install Vagrant
 - <https://www.vagrantup.com/downloads.html>
- Download Vagrant file
 - <https://raw.githubusercontent.com/ogre0403/NCHC-Spark-Tutorial/master/vagrant/Vagrantfile>
- Bring up Lab VMs
 - `vagrant up [spark-master|spark-slave]`
- Shutdown Lab VMs
 - `vagrant halt [spark-master|spark-slave]`

Exercise i-2:

Try interactive Spark shell

- `$cp /opt/spark-1.5.1-bin-hadoop2.6.tgz ~`
- `$tar zxvf spark-1.5.1-bin-hadoop2.6.tgz`
- `$mv spark-1.5.1-bin-hadoop2.6 spark`
- Launch interactive shell
 - `$ ~/spark/bin/spark-shell`
 - `$ ~/spark/bin/pyspark`
- Hadoop-free version v.s. pre-built for Hadoop
 - <http://spark.apache.org/downloads.html>
 - <http://spark.apache.org/docs/latest/hadoop-provided.html>

Shell example

- spark-shell
 - **scala>** val lines = sc.textFile("README.md")
 - **scala>** val pythonLines = lines.filter(s => s.contains("Python"))
 - **scala>** pythonLines.collect()
- pyspark
 - **>>>** lines = sc.textFile("README.md")
 - **>>>** pythonLines = lines.filter(lambda line: "Python" in line)
 - **>>>** pythonLines.collect()

Functional Programming

- Pass FUNCTION (what to do) as method parameter, rather than OBJECT (what)
- Scala: FP language based on JVM
- Java: support FP since Java 8
- Python: support FP

non-Java Lambda v.s. Java 8 Lambda

```
button.addActionListener(new ActionListener() {  
    public void actionPerformed(ActionEvent event) {  
        System.out.println("button clicked");  
    }  
});
```

```
button.addActionListener(event ->  
    System.out.println("button clicked"));
```

```
Thread thread1 = new Thread(new Runnable() {  
    @Override  
    public void run(){  
        System.out.println("Task #1 is running");  
    }  
});
```

```
Runnable task2 = () -> {  
    System.out.println("Task #2 is running"); };
```

```
//Sorting using Anonymous Inner class.  
Collections.sort(personList, new Comparator<Person>(){  
    public int compare(Person p1, Person p2){  
        return p1.firstName.compareTo(p2.firstName);  
    }  
});
```

```
//Anonymous Inner class replaced with Lambda  
expression.  
Collections.sort(personList, (Person p1, Person p2) ->  
    p1.firstName.compareTo(p2.firstName));
```

MapReduce Recap

- The idea of MR originates from FP
- map() and reduce() both are FP operations
- Pass what to do as map() and reduce() parameter

```
public class WordCountMapper
    extends Mapper< Object, Text , Text, IntWritable>{

    public void map(Object key, Text value, Context context )
        throws IOException, InterruptedException {

    }
}
```

Passing function **expression** to Spark

- Using function expression

```
JavaRDD<String> lines = sc.textFile("hdfs://log.txt")  
                        .filter(s -> s.contains("error"));  
long numErrors = lines.count();
```

Passing function **class** to Spark

- We will use function class in our tutorial
 - Override different call() in different Function class

```
Thread thread1 = new Thread(new Runnable() {  
    @Override  
    public void run(){  
        System.out.println("Task #1 is running");  
    }  
});
```

Function name	Method to implement	Usage
Function<T, R>	R call(T)	Take in one input and return one output, for use with operations like map() and filter().
Function2<T1, T2, R>	R call(T1, T2)	Take in two inputs and return one output, for use with operations like aggregate() or fold().

Passing function **class** to Spark

- Using anonymous function class

```
JavaRDD<String> lines = sc.textFile("hdfs://log.txt").filter(  
    new Function<String, Boolean>() {  
        public Boolean call(String s) {  
            return s.contains("error");  
        }  
    });  
long numErrors = lines.count();
```

- Using named class

```
class Contains implements Function<String, Boolean>() {  
    private String query;  
    public Contains(String query) { this.query = query; }  
    public Boolean call(String x) { return x.contains(query); }  
}  
  
RDD<String> errors = lines.filter(new Contains("error"));
```

Exercise i-3:

Pass function to Spark in Java

- Download example code
 - git clone <https://github.com/ogre0403/NCHC-Spark-Tutorial.git>
- 篩選出含 “Python” 的內容
 - `JavaRDD<String> pythonLines = lines.filter(...)`
 - Function class and expression
- 在IDE裡以local 模式執行spark application

Exercise i-4:

Build and launch standalone project


- Scala project
 - Built by sbt, launch by spark-submit
- Python project
 - Launch by spark-submit
- Java project
 - With-Lambda & non-Lambda
 - Built by maven, launch by spark-submit

Outline

- Part I : Quick tour of Spark
 - Why Spark
 - Prepare Lab Environment
 - Run first Spark example
 - {Java / Python / Scala } code via {Shell / Standalone application }
 - Introduction to FP, Java 8, and Scala
- Part II : Programming with RDD
 - RDD Concepts
 - RDD, RDD API, convert RDD
 - Load and save data
 - Spark Application

RDD Essentials

- Resilient distributed dataset
- Each RDD is split into multiple **partitions**
 - Partitions may exist on different machines
- immutable distributed collection of objects
 - Transform creates new RDD
- Spark keeps track lineage graph
 - Fast recovery from failure
- Lazy Evaluation
 - Until action is called



RDD operations

Transformations

- Create a new dataset from and existing one.
- Lazy in nature. They are executed only when some action is performed.
- Example :
 - Map(func)
 - Filter(func)
 - Distinct()

Actions

- Returns to the driver program a value or exports data to a storage system after performing a computation.
- Example:
 - Count()
 - Reduce(func)
 - Collect
 - Take()

Persistence

- For caching datasets in-memory for future operations.
- Option to store on disk or RAM or mixed (Storage Level).
- Example:
 - Persist()
 - Cache()


Create RDD



```
// base RDD  
val lines = sc.textFile("hdfs://...")
```

```
JavaRDD<String> lines = sc.textFile("/path/to/README.md");
```


RDD Transformation



```
// transformed RDDs  
val errors = lines.filter(_.startsWith("ERROR"))  
val messages = errors.map(_.split("\t")).map(r => r(1))  
messages.cache()
```

```
JavaRDD<String> errorsRDD = inputRDD.filter(  
    new Function<String, Boolean>() {  
        public Boolean call(String x) { return x.contains("error"); }  
    }  
));
```

RDD Action




```
// action 1  
messages.filter(_.contains("mysql")).count()
```

```
System.out.println("Input had " + badLinesRDD.count() + " concerning lines")  
System.out.println("Here are 10 examples:")  
for (String line: badLinesRDD.take(10)) {  
    System.out.println(line);  
}
```

Lineage Graph

- Think of each RDD as consisting of instructions on how to compute the data through transformations.

```
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD = inputRDD.filter(lambda x: "warning" in x)
badLinesRDD = errorsRDD.union(warningsRDD)
```



Cache() or Persist()

- Each time call a action, the entire RDD must be computed “from scratch”
 - Cache/persist the computed lineage result

Example 3-39. Double execution in Scala

```
val result = input.map(x => x*x)
println(result.count())
println(result.collect().mkString(","))
```

Example 3-40. persist() in Scala

```
val result = input.map(x => x * x)
result.persist(StorageLevel.DISK_ONLY)
println(result.count())
println(result.collect().mkString(","))
```

- If there are no enough memory for caching RDD partitions
 - MEMORY_ONLY:
 - Oldest will be deleted, then recomputed if necessary
 - MEMORY_AND_DISK:
 - Swap to disk and read back when necessary

Level	Space used	CPU time	In memory	On disk	Comments
MEMORY_ONLY	High	Low	Y	N	
MEMORY_ONLY_SER	Low	High	Y	N	
MEMORY_AND_DISK	High	Medium	Some	Some	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_SER	Low	High	Some	Some	Spills to disk if there is too much data to fit in memory. Stores serialized representation in memory.
DISK_ONLY	Low	High	N	Y	

RDD Type

- Basic RDD[T]
 - considers each data item as a single value
 - Convert to other RDD Type
- PairRDDs
 - each data item containing key/value pairs.
- DoubleRDD
 - Data items are convertible to the Scaladata-type double

<http://blog.csdn.net/pelick/article/details/44922619>

<http://lxw1234.com/archives/2015/07/363.htm>

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>

Create RDD

- 從集合建立RDD
 - `sc.parallelize(1 to 10)`
 - `sc.parallelize(Array("1","2","3"))`
- 從檔案建立RDD
 - `sc.textFile("file://path/to/file")`
 - `sc.textFile("hdfs://path/to/file")`

Basic RDD Transformation (part)

- `RDD.map(func)`

Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>map()</code>	Apply a function to each element in the RDD and return an RDD of the result.	<code>rdd.map(x => x + 1)</code>	<code>{2, 3, 4, 4}</code>

```
JavaRDD<Integer> rdd = sc.parallelize(Arrays.asList(1, 2, 3, 4));
JavaRDD<Integer> result = rdd.map(new Function<Integer, Integer>() {
    public Integer call(Integer x) { return x*x; }
});
System.out.println(StringUtils.join(result.collect(), ","));
```

Basic RDD Transformation (part)

- `RDD.filter(func)`

Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>filter()</code>	Return an RDD consisting of only elements that pass the condition passed to <code>filter()</code> .	<code>rdd.filter(x => x != 1)</code>	<code>{2, 3, 3}</code>

```
RDD<String> errors = lines.filter(new Function<String, Boolean>() {  
    public Boolean call(String x) { return x.contains("error"); }  
});
```

Basic RDD Transformation (part)

- `RDD.flatMap(func)`

Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>flatMap()</code>	Apply a function to each element in the RDD and return an RDD of the contents of the iterators returned. Often used to extract words.	<code>rdd.flatMap(x => x.to(3))</code>	{1, 2, 3, 2, 3, 3, 3}

```
JavaRDD<String> lines = sc.parallelize(Arrays.asList("hello world", "hi"));
JavaRDD<String> words = lines.flatMap(new FlatMapFunction<String, String>() {
    public Iterable<String> call(String line) {
        return Arrays.asList(line.split(" "));
    }
});
words.first(); // returns "hello"
```

Exercise ii-1:

RDD transformation operation

- Create RDD
 - `scala> val rdd = sc.parallelize(1 to 10)`
- `map`
 - `scala> rdd.map(x => x + 1).collect()`
 - `scala> rdd.map(x=>x.toFloat).collect()`
- `flatMap`
 - `scala > rdd.flatMap(x => (1 to x)).collect()`
- `filter`
 - `scala > rdd.filter(x => x > 5)`

Basic RDD Action (part)

- `RDD.collect()`

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>collect()</code>	Return all elements from the RDD.	<code>rdd.collect()</code>	{1, 2, 3, 3}

- `collect()` will attempt to copy every single element in the RDD onto the single driver program, and then run out of memory and crash.

Basic RDD Action (part)

- `RDD.reduce(func)`

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>reduce(func)</code>	Combine the elements of the RDD together in parallel (e.g., sum).	<code>rdd.reduce((x, y) => x + y)</code>	9

```
Integer sum = rdd.reduce(new Function2<Integer, Integer, Integer>() {  
    public Integer call(Integer x, Integer y) { return x + y; }  
});
```


Basic RDD Action (part)

- `RDD.fold(zero)(func)`
 - `Reduce()`的一般式
 - Zero value會運用在每個partition中
 - Zero value也會用在合併partition的結果中

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>fold(zero)(func)</code>	Same as <code>reduce()</code> but with the provided zero value.	<code>rdd.fold(0)((x, y) => x + y)</code>	9

Basic RDD Action (part)

- `RDD.foreach(func)`
 - `Foreach()` iterates over a list and applies some operation with side effects to each list member
 - `Map()` iterates over a list, transforms each member of that list, and returns another list of the same size with the transformed members

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>foreach(func)</code>	Apply the provided function to each element of the RDD.	<code>rdd.foreach(func)</code>	Nothing

Basic RDD Action (part)

- `RDD.saveAsTextFile()`
 - 若路徑為本地路徑，則只會存在executor所在機器上

Exercise ii-2:

RDD action operation

- Reduce
 - `scala> var rdd1 = sc.makeRDD(1 to 10,3)`
 - `scala> rdd1.reduce((a,b) => a+b)`

 - `scala> val range = ('a' to 'z').map(_._toString)`
 - `scala> val rdd = sc.parallelize(range,3)`
 - `scala> rdd.reduce((a,b) => a+b)`
 - `scala > rdd.partitions.size`
- Fold
 - `scala> var rdd1 = sc.makeRDD(1 to 10,3)`
 - `scala> rdd1.fold(0)((a,b) => a+b)`

 - `scala> val range = ('a' to 'z').map(_._toString)`
 - `scala> val rdd = sc.parallelize(range,3)`
 - `scala> rdd.fold("1")((a,b) => a+b)`
 - `scala > rdd.partitions.size`

- `foreach`
 - `scala> var rdd1 = sc.makeRDD(1 to 10,3)`
 - `scala> rdd1.foreach(x=>println(x))`
- `saveAsTextFile`
 - `scala> var rdd1 = sc.makeRDD(1 to 10,3)`
 - `scala> rdd1.saveAsTextFile("/tmp/test")`

RDD Type

- Basic RDD[T]
 - considers each data item as a single value
 - Convert to other RDD Type
- PairRDD[K,V]
 - each data item containing key/value pairs.
- DoubleRDD
 - Data items are convertible to the Scaladata-type double

<http://blog.csdn.net/pelick/article/details/44922619>

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>

Convert between RDD Type

- Scala: Basic RDD convert to
 - PairRDD: Implicit conversion on RDDs of tuple exists.
 - DoubleRDD: Sometimes, implicit convert RDD to RDD[Double]

- Java Basic RDD convert to
 - PairRDD : **explicitly** implement PairFunction()
 - DoubleRDD: **explicitly** implement DoubleFunction()

Function name	Equivalent function*<A, B,...>	Usage
PairFunction<T, K, V>	Function<T, Tuple2<K, V>>	PairRDD<K, V> from a mapToPair
DoubleFunction<T>	Function<T, double>	DoubleRDD from map ToDouble

```
PairFunction<String, String, String> keyData =
    new PairFunction<String, String, String>() {
        public Tuple2<String, String> call(String x) {
            return new Tuple2(x.split(" ")[0], x);
        }
    };
JavaPairRDD<String, String> pairs = lines.mapToPair(keyData);
```

```
JavaDoubleRDD result = rdd.mapToDouble(
    new DoubleFunction<Integer>() {
        public double call(Integer x) {
            return (double) x * x;
        }
    });
System.out.println(result.mean());
```


Exercise ii-3:

RDD conversion

- `scala> val rdd = sc.parallelize(1 to 10)`
- `scala> val pairs = rdd.map(x=> (x,x+1))`
- `scala> pairs.collect()`

- `scala> val rdd1 = sc.parallelize(Array("1","2","3"))`
- `scala> rdd1.mean()`
- `scala> val rdd2 = rdd1.map(x=>x.toInt)`
- `scala> rdd2.mean()`

RDD Type

- Basic RDD[T]
 - considers each data item as a single value
 - Convert to other RDD Type
- PairRDD[K,V]
 - each data item containing key/value pairs.
- DoubleRDD
 - Data items are convertible to the Scaladata-type double

<http://blog.csdn.net/pelick/article/details/44922619>

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>

DoubleRDD

- 針對數值資料(Double)的RDD，提供敘述統計運算
- stats()
 - count()、mean()、sum()、variance()、stdev()、max()、min()

Exercise ii-4: Remove Outliers

- `scala> val rdd1 = sc.parallelize(Array(1,1,1,1,1,1,1,1,1,10))`
- `scala> val distanceDouble = distance.map(string => string.toDouble)`
- `scala> val stats = distanceDoubles.stats()`
- `scala> val stddev = stats.stddev`
- `scala> val mean = stats.mean`
- `scala> val reasonableDistances = distanceDoubles.filter(x => math.abs(x-mean) < 3 * stddev)`

RDD Type


- Basic RDD[T]
 - considers each data item as a single value
- Convert to other RDD Type
- **PairRDD[K,V]**
 - each data item containing key/value pairs.
- DoubleRDD
 - Data items are convertible to the Scaladata-type double

<http://blog.csdn.net/pelick/article/details/44922619>

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>

PairRDD Transformation

- PairRDD is also a RDD
 - RDD.filter(func)
 - RDD.map(func)
 - RDD.flatMap(func)
 - ...



```
Function<Tuple2<String, String>, Boolean> longWordFilter =  
    new Function<Tuple2<String, String>, Boolean>() {  
        public Boolean call(Tuple2<String, String> keyValue) {  
            return (keyValue._2().length() < 20);  
        }  
    };  
JavaPairRDD<String, String> result = pairs.filter(longWordFilter);
```

PairRDD Transformation

- `RDD.mapValues(func)`

Table 4-1. Transformations on one pair RDD (example: {(1, 2), (3, 4), (3, 6)})

Function name	Purpose	Example	Result
<code>mapValues(func)</code>	Apply a function to each value of a pair RDD without changing the key.	<code>rdd.mapValues(x => x+1)</code>	{(1, 3), (3, 5), (3, 7)}

```
JavaPairRDD<Integer,Integer> result =  
    prdd.mapValues(new Function<Integer, Integer>() {  
        @Override  
        public Integer call(Integer v1) throws Exception {  
            return v1 +1;  
        }  
    });
```

PairRDD Transformation

- `RDD.reduceByKey(func)`




Table 4-1. Transformations on one pair RDD (example: $\{(1, 2), (3, 4), (3, 6)\}$)

Function name	Purpose	Example	Result
<code>reduceByKey(func)</code>	Combine values with the same key.	<code>rdd.reduceByKey((x, y) => x + y)</code>	$\{(1, 2), (3, 10)\}$

```
JavaPairRDD<Integer,Integer> result =
    prdd.reduceByKey(new Function2<Integer, Integer, Integer>() {
        @Override
        public Integer call(Integer v1, Integer v2) throws Exception {
            return v1 + v2;
        }
    });
```


PairRDD Transformation

- `RDD.groupByKey()`

GroupByKey


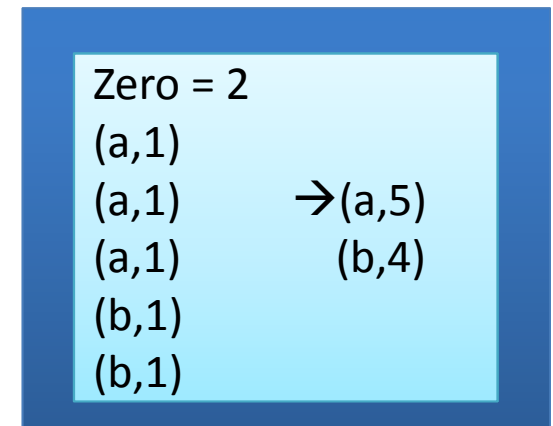




Table 4-1. Transformations on one pair RDD (example: {(1, 2), (3, 4), (3, 6)})

Function name	Purpose	Example	Result
<code>groupByKey()</code>	Group values with the same key.	<code>rdd.groupByKey()</code>	<code>{{(1, [2]), (3, [4, 6])}}</code>

PairRDD Transformation

- `RDD.foldByKey(zero)(func)`
 - `reduceByKey()`的一般式
 - Zero value會運用在每個partition同key的值



PairRDD Action

- `RDD.countByKey()`

Table 4-3. Actions on pair RDDs (example $\{(1, 2), (3, 4), (3, 6)\}$)

Function	Description	Example	Result
<code>countByKey()</code>	Count the number of elements for each key.	<code>rdd.countByKey()</code>	$\{(1, 1), (3, 2)\}$

Exercise ii-5:

Pair RDD

- `scala> var rdd1 =
sc.makeRDD(Array(("A",0),("A",2),("B",1),("B",
2),("C",1)))`
- `scala > rdd1.reduceByKey(_ + _).collect`
- `scala > rdd1.groupByKey()`
- `scala > rdd1.countByKey()`
- `scala> rdd1.foldByKey(0)(_+_).collect`
- `scala> rdd1.foldByKey(2)(_+_).collect`


- `scala> var rdd1 =
sc.makeRDD(Array(("A","a"),("A","b"),("A","c")
,("B","a"),("B","b")),3)`
- `scala> rdd1.foldByKey("0")(_+_).collect`
- `scala> rdd1.mapPartitionsWithIndex((index:
Int, it: Iterator[(String,String)])
=>it.toList.map(x => if (index ==0)
{println(x)}).iterator).collect`

Spark application

- Pi Estimation
- Word Count
- K-means
- Apache Log analysis

Pi Estimation

- 在單位方型內隨機生成資料點。
 - `mapToPair(...)`
- 判斷那些點在單位圓內，那些點在單位圓外面。
 - `filter(...)`
- 統計在單位圓內點的個數
 - `count()`
- $\text{Pi} = 4 * \text{圓內點總數} / \text{全部點數}$



$$\frac{\text{Area of Circle}}{\text{Area of Square}} = \frac{\pi r^2}{(2r)^2} = \frac{\pi}{4}$$

Exercise ii-6: Pi Estimation

- 分別用 Function class 和 expression 完成
 - mapToPair(...)
 - filter(...)

MR convert to Spark

- Map() in MR
 - flatmap() + map()/mapToPair()
- Reduce() in MR
 - reduceByKey()
 - foldByKey() (reduceByKey() 的一般化)
 - groupByKey() + mapValue() (效能差，不建議使用)

```
JavaRDD<String>
-----
aaa bbb ccc
bbb ccc ddd
```



flatMap



```
JavaRDD<String>
-----
aaa
bbb
ccc
bbb
ccc
ddd
```

```
JavaRDD<String>
-----
aaa
bbb
ccc
bbb
ccc
ddd
```



mapToPair



```
JavaPairRDD<String,int>
-----
aaa      1
bbb      1
ccc      1
bbb      1
ccc      1
ddd      1
```

Key	value
aaa	1

Key	value
bbb	1
Bbb	1

Key	value
ccc	1
ccc	1

Key	value
ddd	1



reducerByKey




Key	value
aaa	1
bbb	2
ccc	2
ddd	1

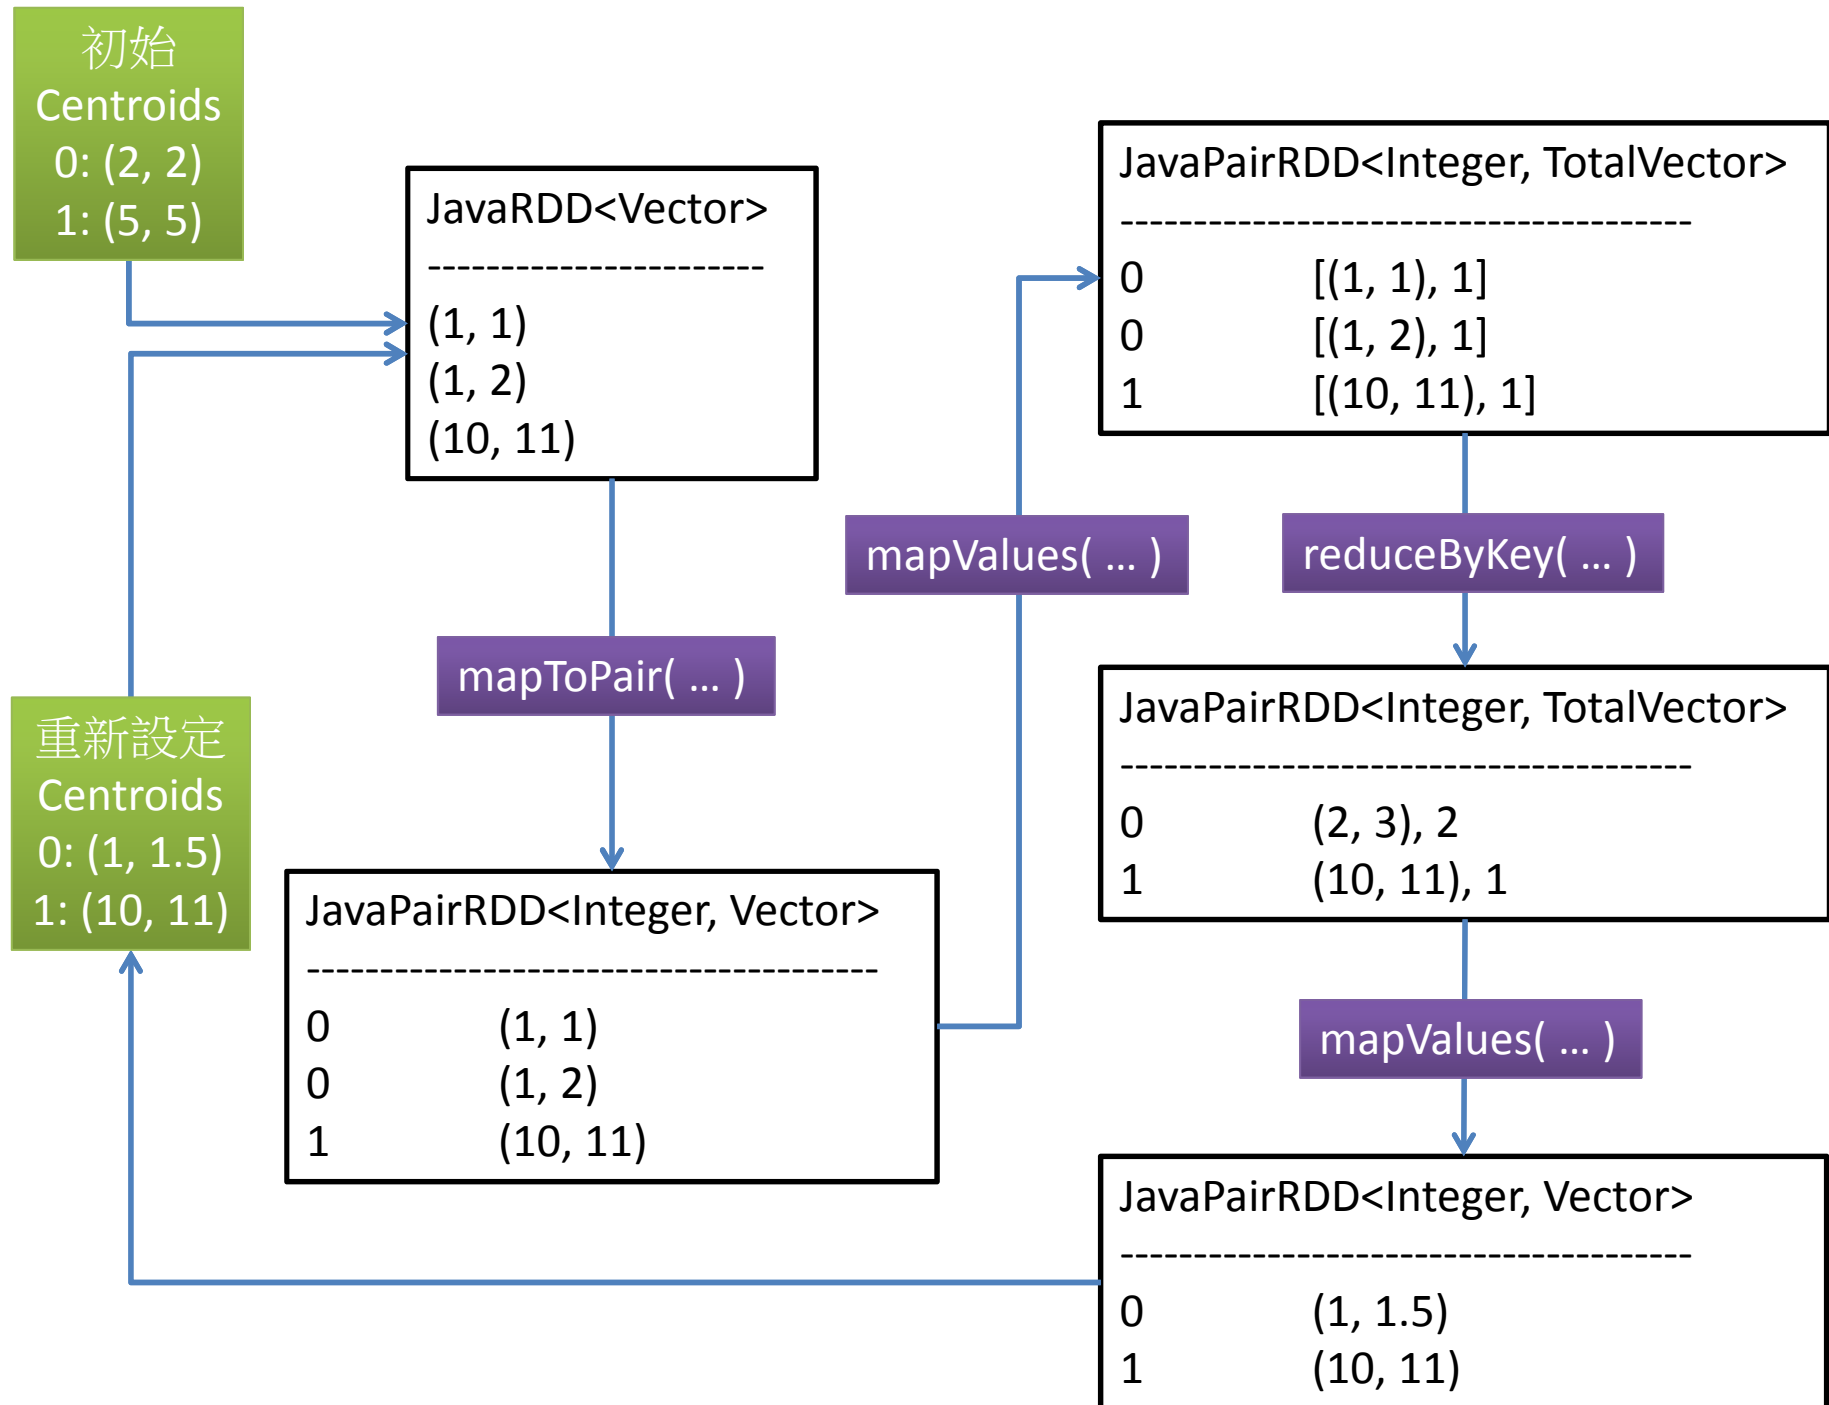
Exercise ii-7: Word Count

- 用 `foldByKey(...)` 改寫 word count
- 用 `groupByKey() + mapValue(...)` 改寫 word count

K-means

- 隨機選取資料組中的k筆資料當作初始群中心 $u_1 \sim u_k$
- 計算每個資料 x_i 對應到最短距離的群中心 (固定 u_i 求解所屬群 S_i)
- 利用目前得到的分類重新計算群中心 (固定 S_i 求解群中心 u_i)
- 重複step 2,3直到收斂 (達到最大疊代次數 or 群心中移動距離很小)





Exercise ii-8:

K-means

- 分群, 計算每個Vector離那一組中心最近
 - `JavaPairRDD<Integer, Vector> result1 = data1.mapToPair(...);`
- 將PairRDD的Vector value轉成TotalVector方便後續做reduceByKey
 - `JavaPairRDD<Integer, TotalVector> result2 = result1.mapValues(...);`
- 將所有點的座標加總
 - `JavaPairRDD<Integer, TotalVector> result3 = result2.reduceByKey(...);`
- 計算所有點的幾何中心, 找出新的centroid 座標
 - `JavaPairRDD<Integer, Vector> result4 = result3.mapValues(...);`

Exercise ii-9:

Apache Log analysis

- Format
 - 64.242.88.10 - - [07/Mar/2004:16:47:12 -0800] "GET /robots.txt HTTP/1.1" 200 68
 - 1:IP 2:client 3:user 4:date time 5:method 6:req 7:proto 8:respcode 9:size
- 用 `parseFromLogLine()` 將日志字符串轉成 `ApacheAccessLog` 物件
 - `JavaRDD<ApacheAccessLog> result1 = logLines.map(...);`
- 將不符合格式的資料篩除
 - `JavaRDD<ApacheAccessLog> result2 = result1.filter(...);`
- 利用 `ApacheAccessLog` 的 `getContentSizeDouble()`, 產生 `DoubleRDD` 做敘述統計
 - `JavaDoubleRDD contentSizes = accessLogs.mapToDouble(...).cache()`

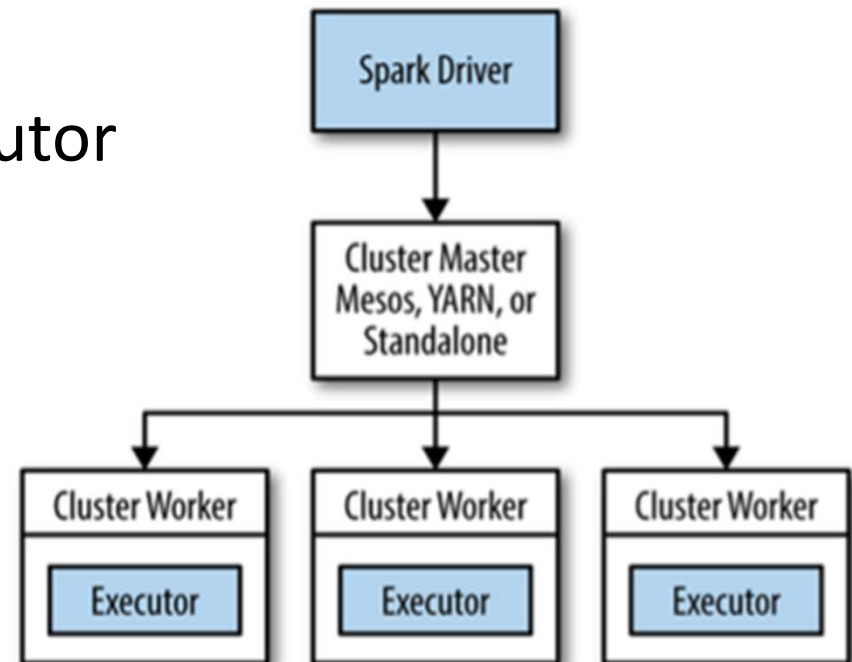
- `countByKey()`找出各種回傳值總數
 - `Map<Integer, Object> responseCodeToCount = accessLogs.mapToPair(...).countByKey()`
- 篩選出大於threshold的ip
 - `JavaPairRDD<String, Long> result6 = result5.filter(...);`
- 利用 `map()` 將 `JavaPairRDD<String, Long>` 轉成 `JavaRDD<String>`
 - `JavaRDD<String> result7 = result6.map(...);`
- 印出RDD中的每個元素
 - `result7.foreach(...)`
- 利用 `mapValues()` 計算 `groupByKey()` 結果的總合
 - `JavaPairRDD<String, Long> result9 = result8.groupByKey().mapValues(...);`

Outline

- Part III: Run on Cluster
 - Cluster Manager
 - Standalone, YARN, Mesos
 - Components of Execution
 - Application, Job, stage, task, RDD, partition, DAG
- Part IV: Play with Spark Ecosystem
 - Spark SQL
 - Spark Streaming
 - SparkR

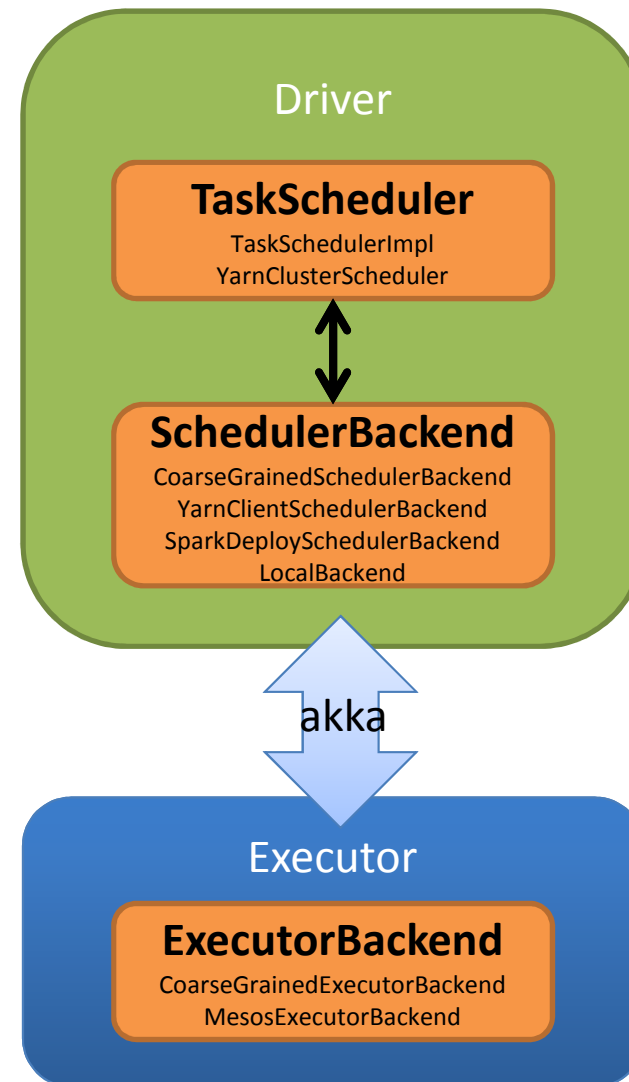
Spark Runtime Architecture

- Driver
 - Process which has the main() method
 - Convert application to tasks
 - DAGScheduler
 - Scheduling tasks on executor
 - TaskScheduler
 - SchedulerBackend
- Executor
 - Running individual task



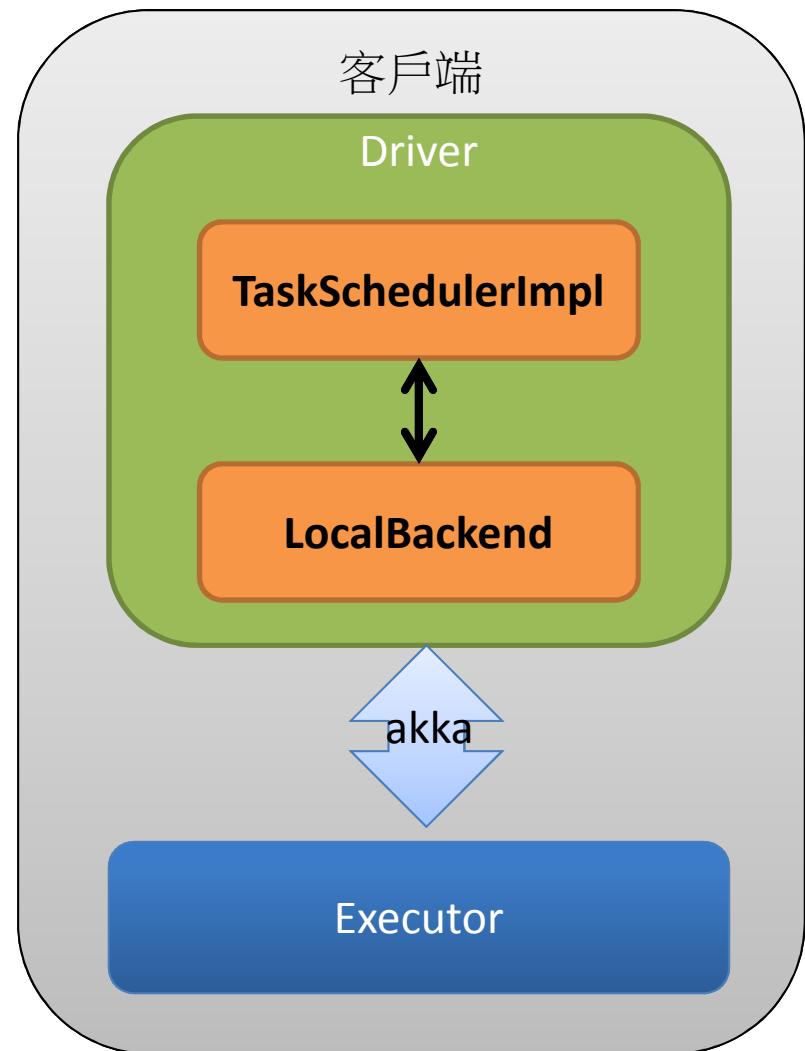
Spark Deployment

- Spark on Local
- Spark on Cluster
 - Spark Standalone
 - Cluster / Client
 - Spark on YARN
 - YARN-Cluster / YARN-Client
 - Spark on Mesos
 - Coarse grained
 - Fine grained



Local mode

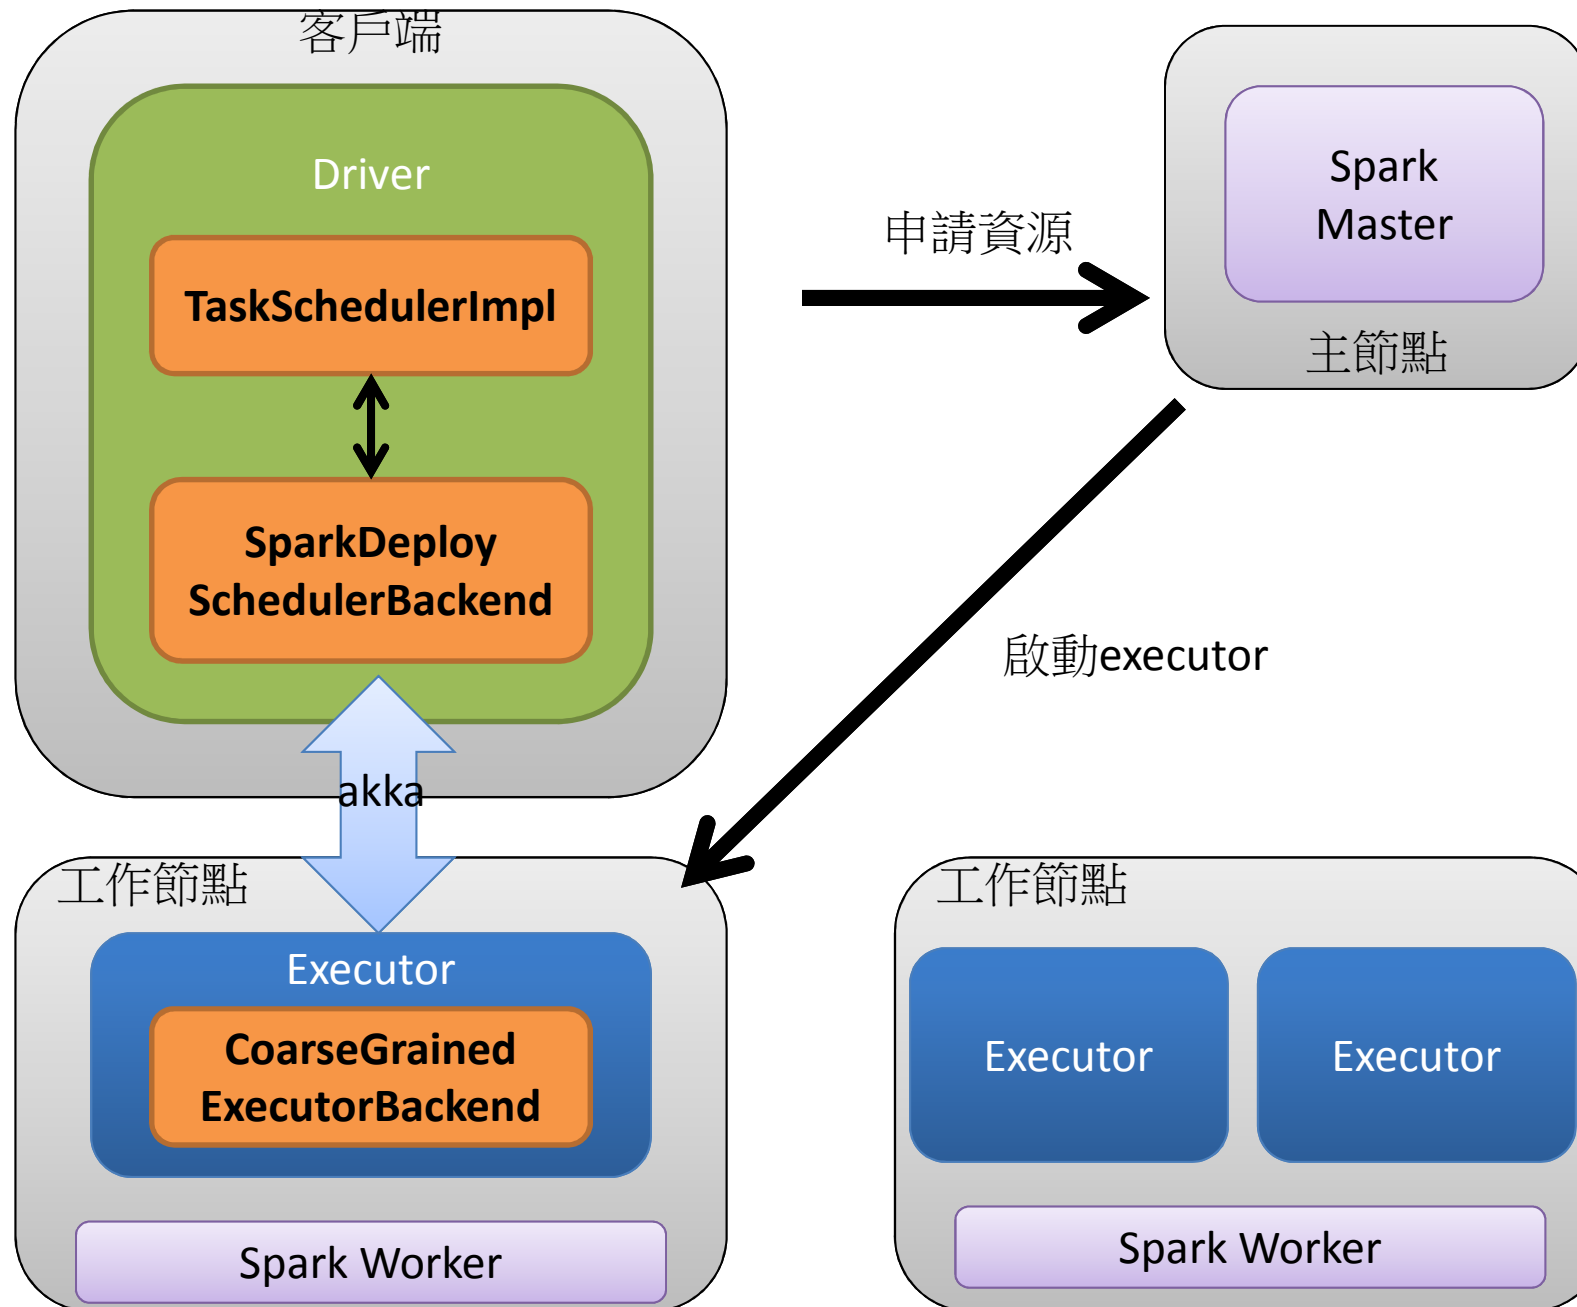
- local
 - local mode with 1 core
- local[N]
 - local mode with N core
- local[*]
 - Local mode with as many cores the node has



Exercise iii-1: use different mode

- `$./spark-shell --master local`
 - **scala>** `sc.master`
 - **scala>** `sc.isLocal`
- `$./spark-submit --master local[2] --class xxx.yyy.zzz xyz.jar`
 - **DON'T** hard code `setMaster()` in source

Standalone mode



Exercise iii-2:

Setup Standalone cluster

- Step 1: edit /etc/hosts
- Step 2: Password less login
 - spark@master \$ ssh-keygen -t rsa
 - spark@master \$ ssh-copy-id spark@spark-master
 - spark@master \$ ssh-copy-id spark@spark-slave
- Step 3:
 - Edit \$SPARK_HOME/conf/slaves
 - Edit \$SPARK_HOME/conf/spark-env.sh (OPTION)
 - SPARK_WORKER_MEMORY=512m
- Step 4: distribute \$SPARK_HOME to all nodes
 - \$scp -r sprak spark@spark-slave:/home/spark

- Step 5: Start / Stop Spark Cluster
 - \$SPARK_HOME/sbin/start-all.sh
 - \$SPARK_HOME/sbin/stop-all.sh
- Step 6: connect to Spark master using **hostname**
 - \$ spark-shell --master spark://**spark-master**:7077
 - \$ spark-submit --master spark://**spark-master**:7077
--class xx.yy.zz xyz.jar
 - export MASTER=spark://spark-master:7077
- Step 7: submit using cluster mode
 - \$ spark-submit **--deploy-mode cluster** \
--master XXX --class xx.yy.zz \
--executor-memory YYM --driver-memory ZZM \
xyz.jar

- Step 8: Setup HDFS (OPTION)
 - Edit \$HADOOP_HOME/libexec/hadoop-config.sh
 - export JAVA_HOME=/usr/lib/jvm/java-8-oracle
 - Edit \$HADOOP_HOME/etc/hadoop/hadoop-env.sh
 - export JAVA_HOME=/usr/lib/jvm/java-8-oracle
 - Edit \$HADOOP_HOME/etc/hadoop/slaves
 - Edit \$HADOOP_HOME/etc/hadoop/core-site.xml
 - Edit \$HADOOP_HOME/etc/hadoop/hdfs-site.xml
 - Edit ~/.bashrc
 - Start / stop HDFS:
 - start-dfs.sh / stop-dfs.sh
 - Read file from HDFS:
 - sc.textFile("hdfs://spark-master:9000/path/to/file")

core-site.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://spark-master:9000</value>
  </property>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/home/spark/hadoop_dir</value>
  </property>
</configuration>
```

hdfs-site.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.permissions</name>
    <value>>false</value>
  </property>
</configuration>
```

```
export HADOOP_HOME=/home/spark/hadoop
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
```

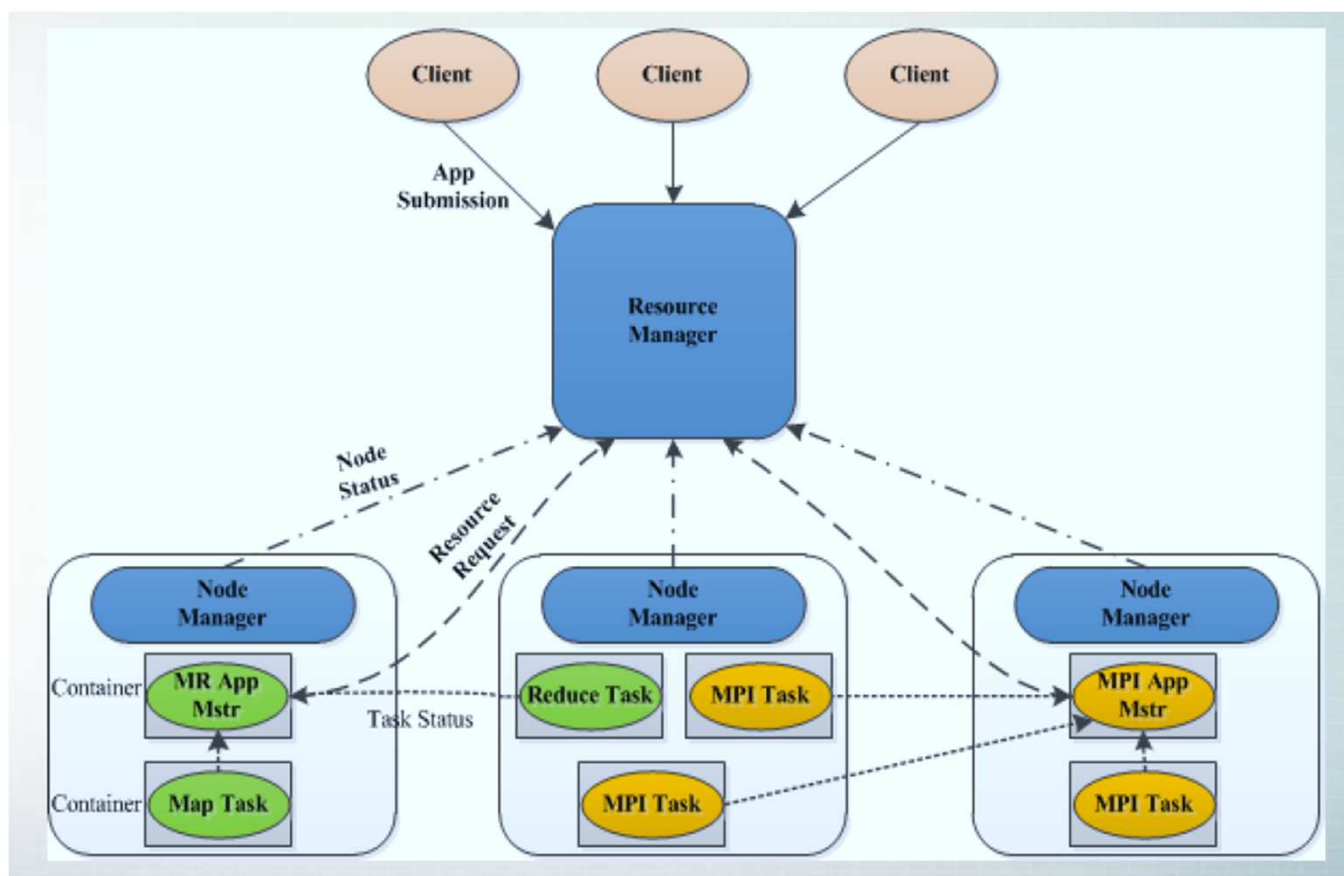
```
export YARN_HOME=$HADOOP_HOME
export YARN_CONF_DIR=$HADOOP_HOME/etc/hadoop
```

.bashrc

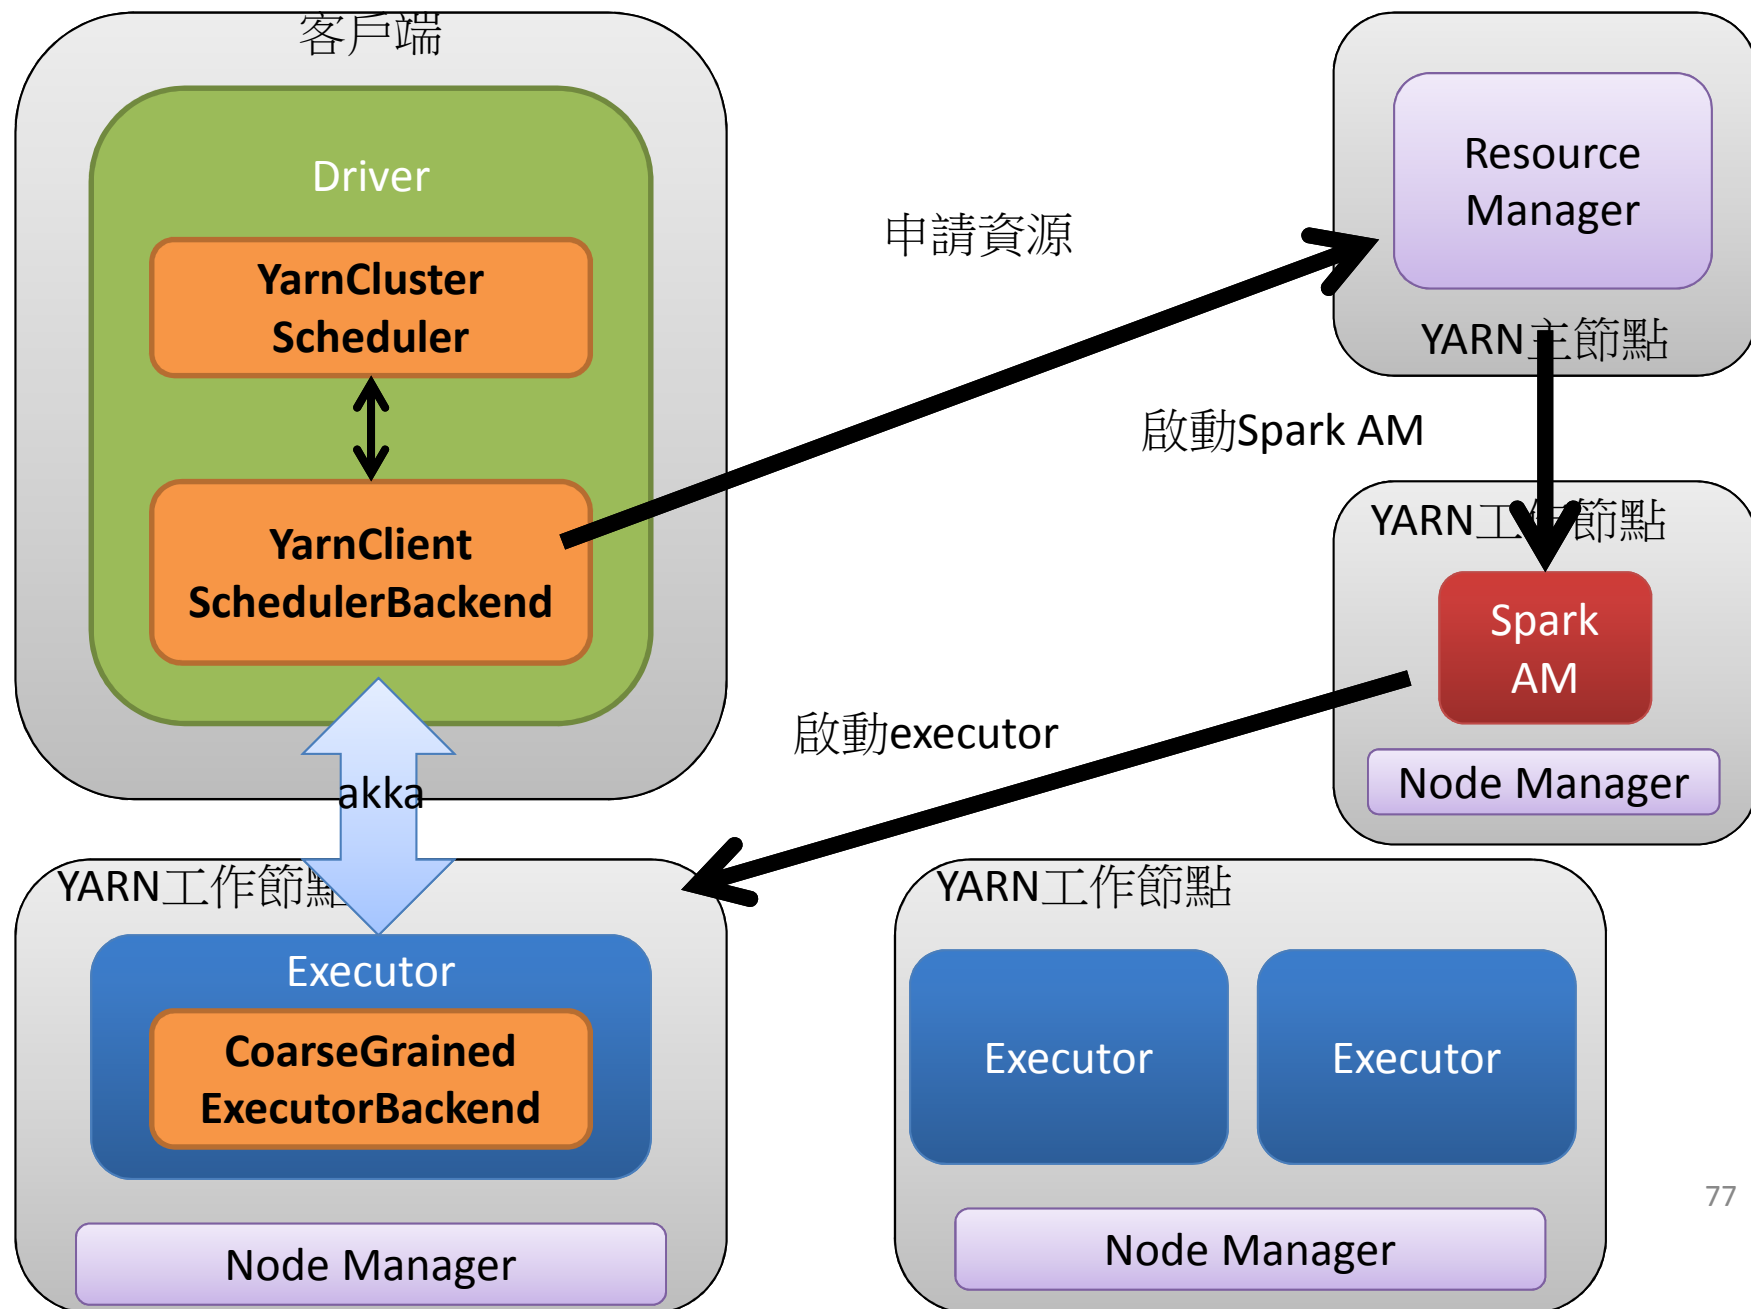
```
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
```

Spark on YARN

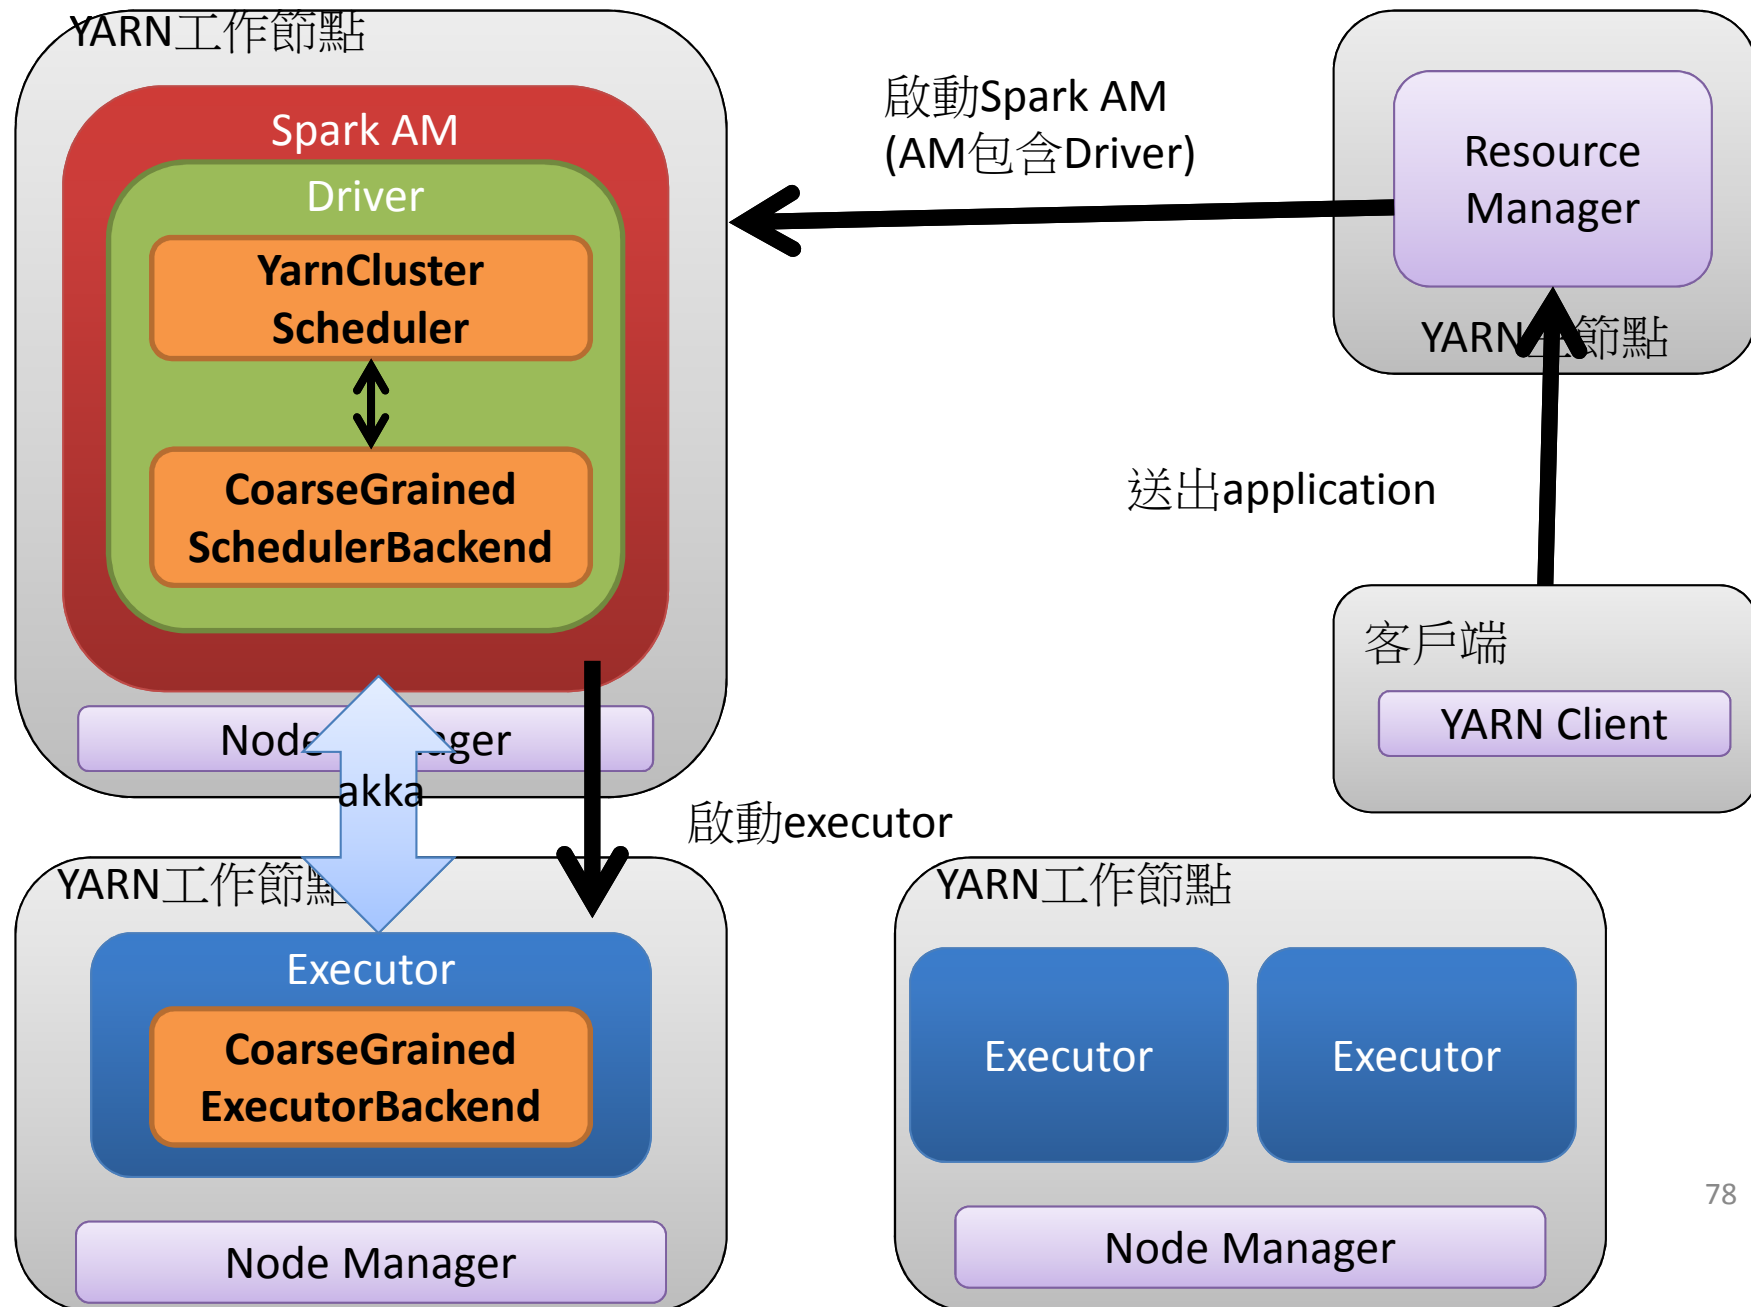
1. 由RM做全局的資源分配
 2. NM定時回報目前的資源使用量
 3. 每個JOB會有一個負責的AppMaster控制Job
 4. 將資源管理與工作控制分開
 5. YARN為一通用的資源管理系統
- 可達成在YARN上運行多種框架



YARN Client mode



YARN Cluster mode



Exercise iii-3:

Setup Spark on YARN cluster

- Step 1: Start HDFS
- Step 2: Configure YARN
 - Edit \$HADOOP_HOME/etc/hadoop/mapred-site.xml
 - Edit \$HADOOP_HOME/etc/hadoop/yarn-site.xml
 - Distribute \$HADOOP_HOME to all nodes
 - \$SPARK_HOME **NO NEEDED** to distribute
- Step 3: Start / Stop YARN Cluster
 - start-yarn.sh / stop-yarn.sh

- Step 4: connect to YARN cluster using **yarn**
 - \$ spark-shell --master yarn
 - \$ spark-submit **--master yarn** \
--class xx.yy.zz xyz.jar
- Step 5: using yarn cluster mode
 - \$ spark-submit --master yarn \
--deploy-mode cluster \
--class xx.yy.zz xyz.jar

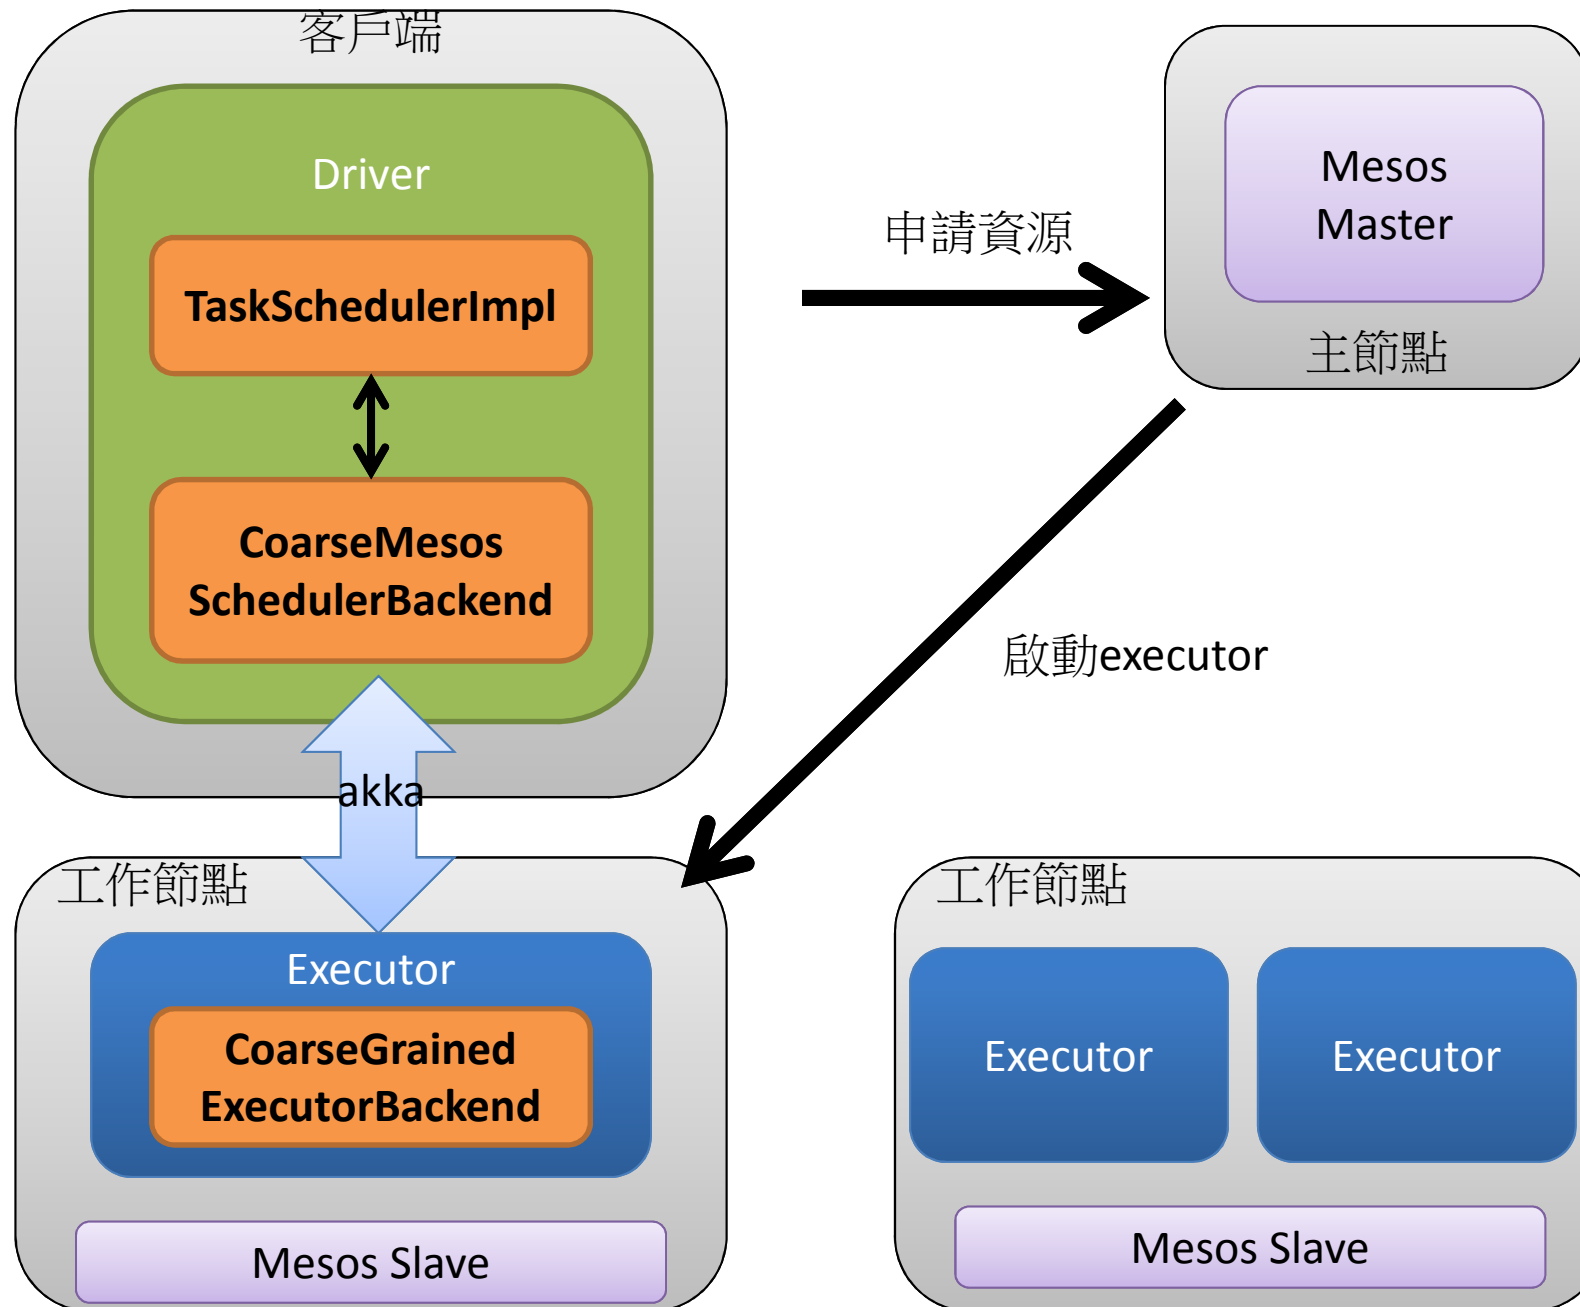
yarn-site.xml

```
<?xml version="1.0"?>
<configuration>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>spark-master:8030</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>spark-master:8031</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>spark-master:8032</value>
  </property>
  <property>
    <name>yarn.nodemanager.address</name>
    <value>0.0.0.0:8034</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.local-dirs</name>
    <value>/home/spark/hadoop_dir/nm-local-dir</value>
  </property>
  <property>
    <name>yarn.nodemanager.log-dirs</name>
    <value>/home/spark/hadoop_dir/userlogs</value>
  </property>
</configuration>
```

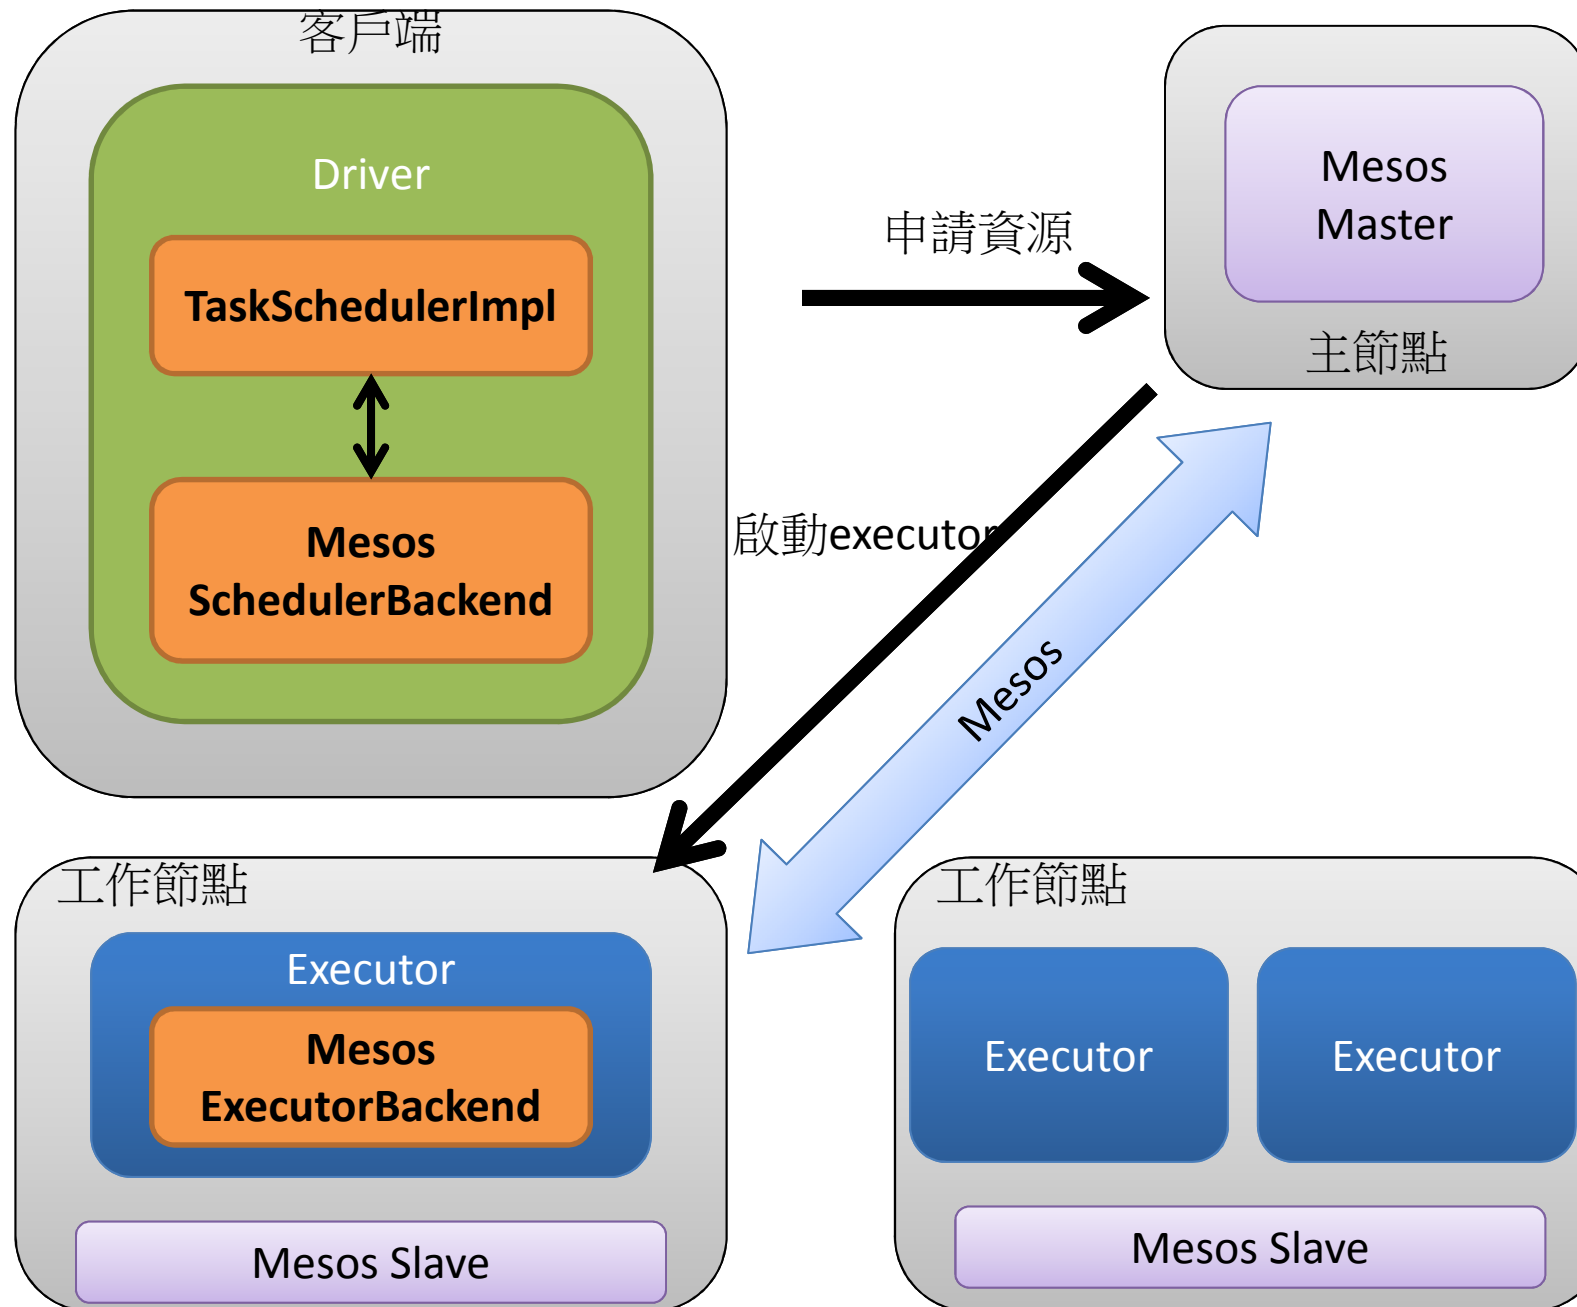
mapred-site.xml

```
<?xml version="1.0"?>
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

Mesos Coarse Grained



Mesos Fine Grained



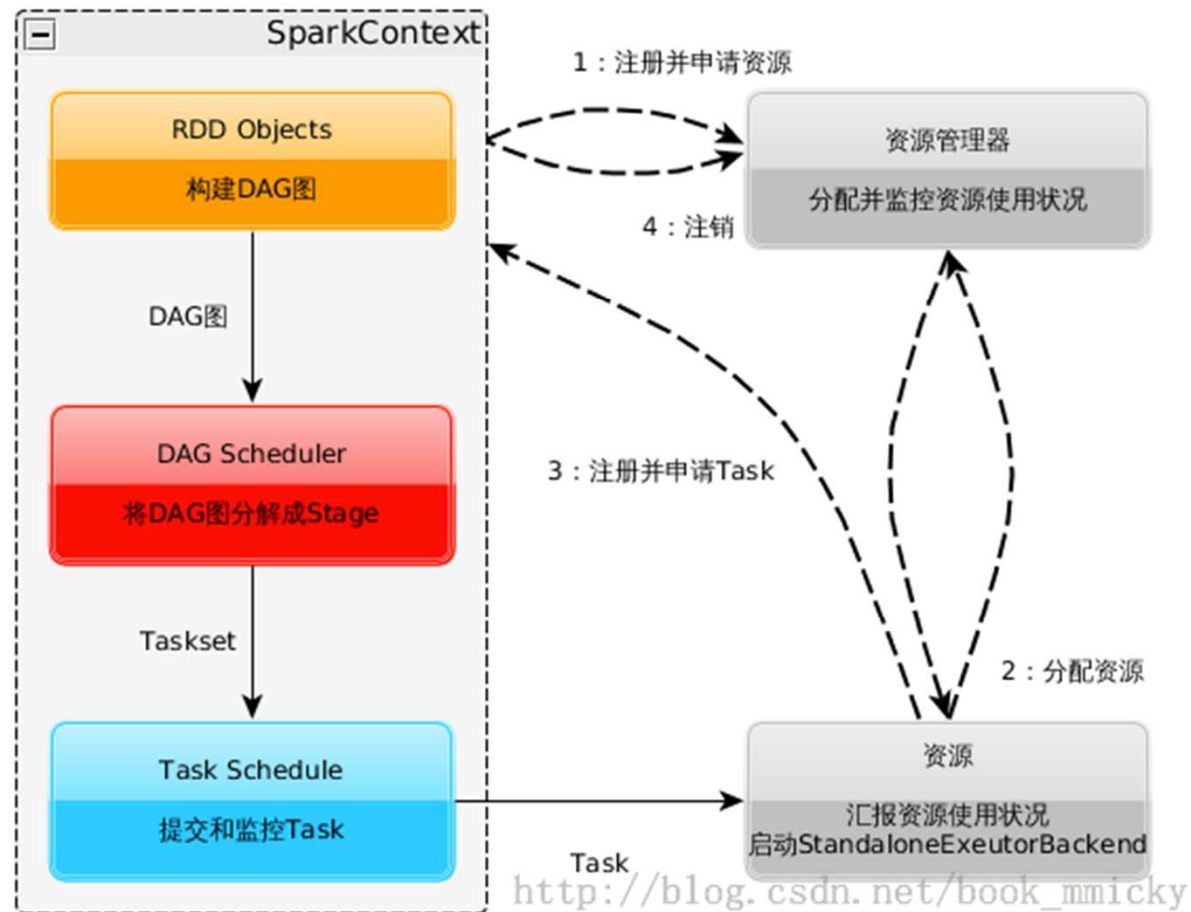
選擇那一種Cluster mode

- 初學者、沒有特殊需求與包袱
 - Standalone mode
- 同時要跑Spark與MapReduce
 - Spark on YARN
- 需要fine grained scheduling
 - Spark on Mesos

Outline

- Part III: Run on Cluster
 - Cluster Manager
 - Standalone, YARN, Mesos
 - Components of Execution
 - Application, Job, stage, task, RDD, partition, DAG
- Part IV: Play with Spark Ecosystem
 - Spark SQL
 - Spark Streaming
 - SparkR

Spark Runtime



- 建構Spark Application執行環境(啟動sc)
- SC透過cluster manager (Standalone、Mesos、Yarn)申請Executor資源後啟動ExecutorBackend，executor會向sc註冊。
- **SC利用 DAGScheduler生成 DAG 圖(logic plan)，將 DAG 圖分解成 Stage(physical plan)，將 Taskset發送給Task Scheduler，最後由Task Scheduler將Task分發給Executor執行。**
- Task在Executor上執行完成後，釋放資源。

Hadoop application

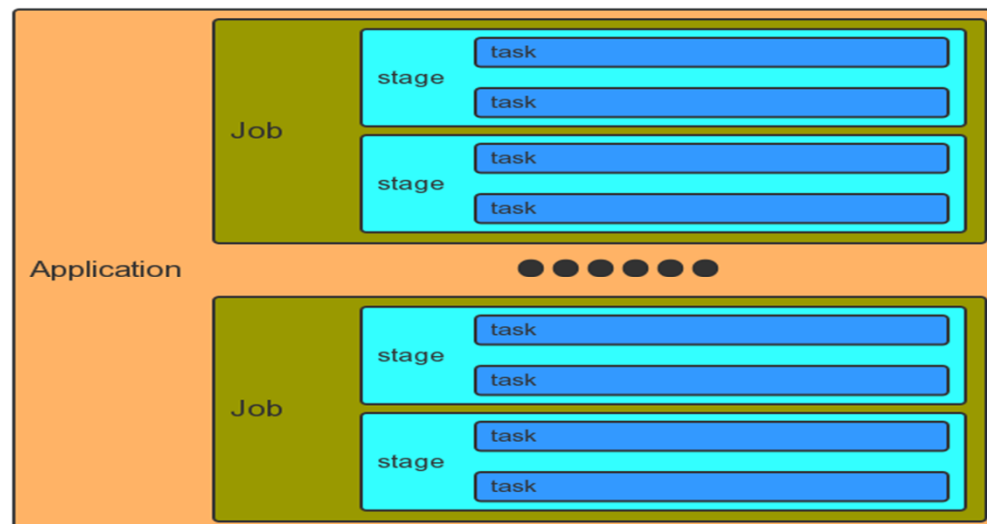
- 一個Mapreduce application就是一個Job
- 一個Job裡會有多個task
- Task 可分成map task與reduce task



Spark技术博客: <http://www.iteblog.com>

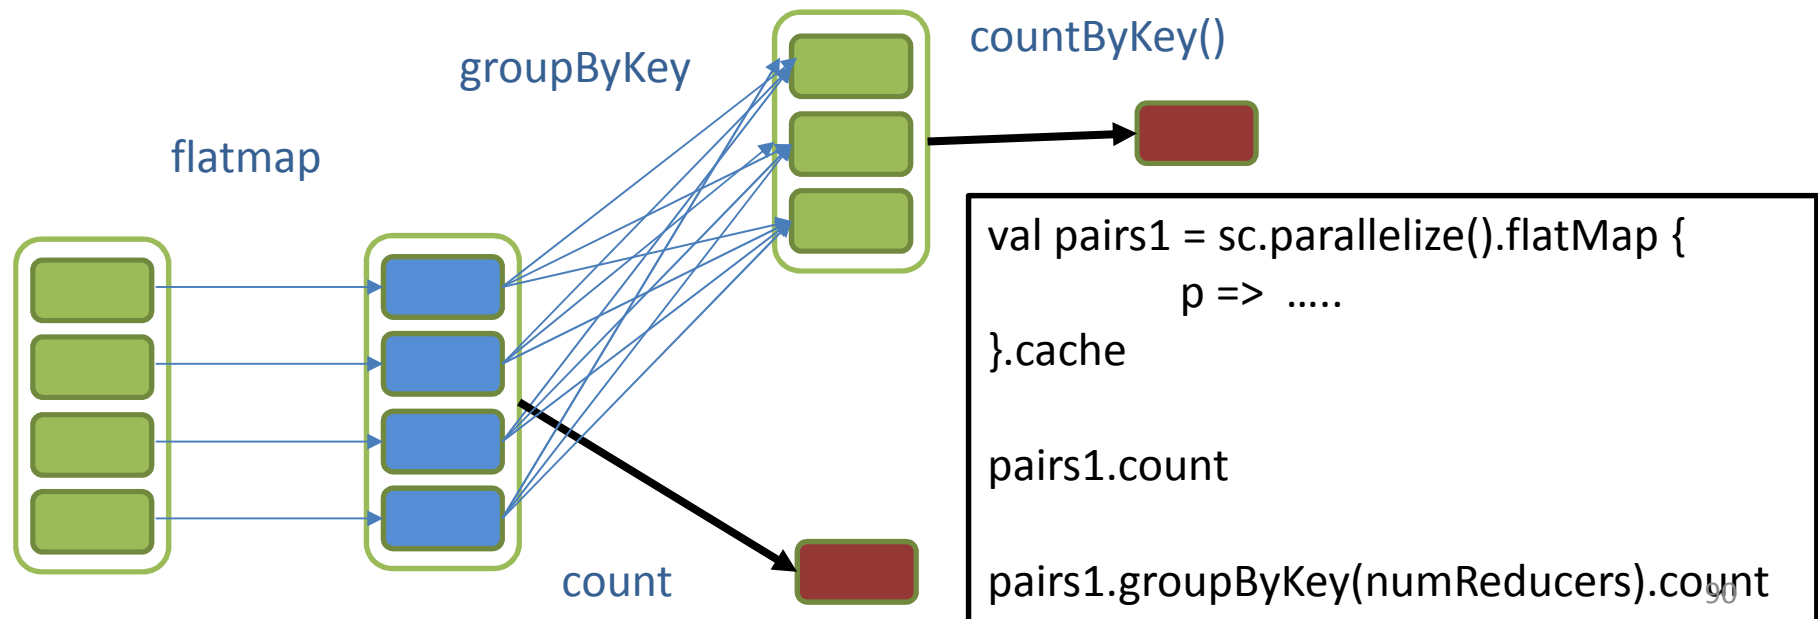
Spark Application

- 一個sparkcontext對應一個Application
- 每個applicatio可有多個job，一個action產生一個job，job可平行或依序執行
- Job由多個stage組成，job裡的stage是由shuffle畫分
- Stage裡有多個task，task數由最後一個RDD的partition決定




How to determine Job

- application 可以包含多個job
- Driver 中每執行一個action，就會產生一個job




Narrow Dependency

- depends only on data that is already residing in the partition and data shuffling is unnecessary
- Filter() 、 map() 、 ...




Narrow transformation

- Input and output stays in same partition
- No data movement is needed




Wide Dependency

- depends on data residing in multiple partitions and therefore data shuffling is needed to bring them together in one place
- `groupByKey()` 、 `reduceByKey()`



Wide transformation

- Input from other partitions are required
- Data shuffling is needed before processing



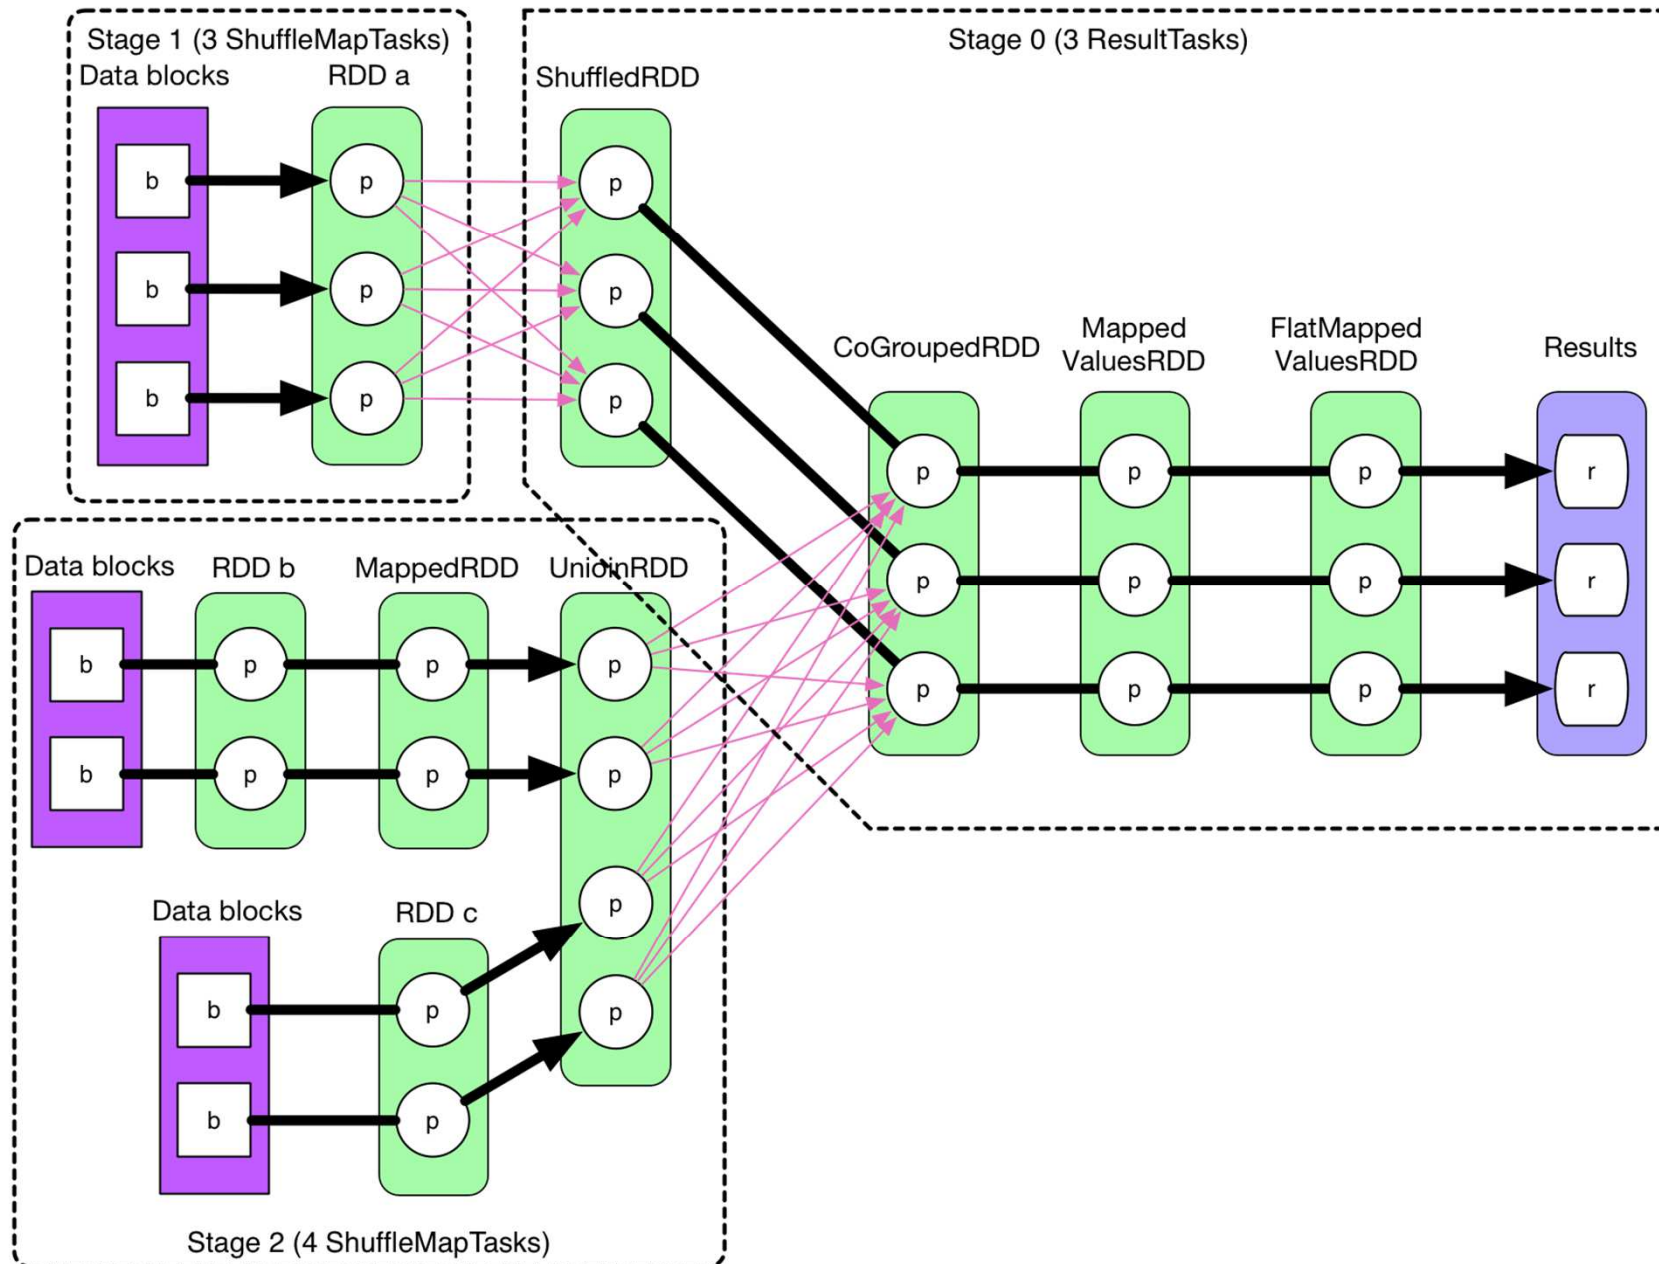
<http://spark-internals.books.yourtion.com/markdown/2-JobLogicalPlan.html>

<http://horicky.blogspot.tw/2015/02/big-data-processing-in-spark.html>

How to determine stage and task

- Stage
 - Stages are sequences of RDDs, that don't have a **Shuffle** in between
 - 由後往前推，遇到wideDependency (or Shuffle) 就斷開，遇到NarrowDependency 就加入stage。
- Task
 - 每個Stage裡的 task數目由該 stage最後一個 RDD 中的 partition個數決定。

ComplexJob
including map(), partitionBy(), union(), and join()



Exercise iii-4 :

Setup history server

- Start HDFS and create log directory
- Edit `$SPARK_HOME/conf/spark-defaults.conf`
 - 讓Executor 輸出 event log
 - `spark.eventLog.enabled` `true`
 - `spark.eventLog.dir` `hdfs://spark-master:9000/sparkLogs`
 - 設定 history server 讀取目錄
 - `spark.history.fs.logDirectory` `hdfs://spark-master:9000/sparkLogs`
- Distribute `spark-defaults.conf` to all nodes
- Start history server
 - `$SPARK_HOME/sbin/start-history-server.sh`

Exercise iii-5 :

觀察Job、Stage、Task in UI

- 分別執行simpleJob與complexJob
- 在job-history-server上驗證job數、stage數、task數

Reference

- Spark
 - Learning Spark
 - Advanced Analytics with Spark
- Java 8
 - Java 8 Lambdas
- Scala
 - Scala in action
 - 为Java程序员编写的Scala的入门教程
 - <http://www.iteblog.com/archives/1325>

